

## Lecture 1. Text Pre-processing

### Definations:

- **Words:** Sequence of characters with a meaning and/or function
- **Sentence:** "The student is enrolled at the University of Melbourne."
- **Word token:** each instance of "the" in the sentence above.
- **Word type:** the distinct word "the".
- **Lexicon:** a group of word types.
- **Document:** one or more sentences.
- **Corpus:** a collection of documents.

### Text Normalisation

- Remove unwanted formatting (e.g.HTML)
- Segment structure(e.g.sentences)
- Tokenise words
- Normalise words
- Remove unwanted words

### Word normalisation

- Lowercasing(Australia->australia)
- Removing morphology (Inflectional and Derivational)
- Correcting spelling
- Expanding abbreviations

**Inflectional Morphology** -> grammatical variants

**Derivational morphology** -> distinct words (suffixes and prefixes)

### Lemmatisation:

Removing any inflection to reach the uninflected form, the lemma.

### Stemming:

Stemming strips off all suffixes, leaving a stem.

## Lecture 2 Information Retrieval: the Vector space model

### Evaluation on Test collections:

IR research often use reusable test collections constructed for reproducible IR evaluation.

### Document representation

Bag-of-Word (BOW) -> term-document matrix

### Cosine distance

$$\text{Cos}(a, b) = \frac{a \cdot b}{|a| \times |b|}$$

### TF-IDF

TF: The number of the term occurs in a specific document.

IDF: The number of document that the term occurs.

### Query processing in VSM

Normalise TF-IDF square root -> divide each item of rows

Similarity: Doc dot Query

### BM25

$$w_{td} = \left[ \log \frac{N - f_t + 0.5}{f_t + 0.5} \right] \quad (\text{idf})$$
$$\times \frac{(k_1 + 1) f_{d,t}}{k_1 \left( (1 - b) + b \frac{L_d}{L_{ave}} \right) + f_{d,t}} \quad (\text{tf and doc. length})$$
$$\times \frac{(k_3 + 1) f_{q,t}}{k_3 + f_{q,t}} \quad (\text{query tf})$$

### Term-wise processing

With the query as a pseudo-document need only dot-product consider terms that are in the query and in the document.

### Inverted Index

- Terms as rows
- Values as lists of (docID, weight) pairs, aka posting list
- (weights listed are the normalised TF\*IDF values)

### Lecture 3: Index compression and efficient query processing

**Entropy:** Information content of a text T

$$H(T) = - \sum_{s \in \Sigma} \frac{f_s}{n} \log_2 \frac{f_s}{n}$$

where  $f_s$  is the frequency of symbol  $s$  in  $T$  and  $n$  is the length of  $T$ .

**Gaps:**

Gaps between two adjacent integers can be much smaller

**Variable Byte Compression:**

Use variable number of bytes to represent integers. Each byte contains 7 bits “payload” and one continuation bit.

**Fast Searching**

1. Binary search over uncompressed sample values to find destination block
2. Decompress destination block to determine final offset in postings list

**Top-k Retrieval**

Retrieve the top  $k$  items for a given query without having to evaluate all documents.

**The WAND Algorithm**

- Keep track of the top- $k$  highest scored documents.
- For each unique term in the collection store the maximum contribution it can have to any document score in the collection.
- Skip over documents that can not enter the top- $k$  highest results.

**Maximum contribution**

The Maximum contribution of a term  $q$  as the largest score any document in the collection can have for the query  $Q$  only consisting of  $q$ .

## Lecture 4: Query Completion and Expansion

### High Level Algorithm:

1. Retrieve set of candidates
2. Rank candidates by frequency
3. Re-rank highest ranked candidates and return top-k

### Completion Types:

- Prefix match.
- Substring match.
- Multi-term prefix match.
- Relaxed match.

### Prefix match (Trie+RMQ based Index)

Store array with frequencies corresponding to each query. Subtree corresponds to range in frequency array. Find the top-K highest numbers in that range.

### Range Maximum Queries

Given an array  $A$  of  $n$  numbers, and a range  $[l, r]$  of size  $m$ , find the positions of the  $K$  largest numbers in  $A[l, r]$ .

1. Find position of largest element of  $A[i, j]$ .
2. Recurse to  $A[i, p-1]$  and  $A[p+1, j]$ .
3. Keep going until you have the  $K$  largest elements.
4. Runtime  $O(K \log K)$ .
5. Instead of precomputing all  $O(n^2)$  ranges  $A[i, j]$ , for each position  $A[i]$ , precompute only  $\log n$  ranges of increasing size:  $A[i, i+1], A[i, i+2], A[i, i+4], A[i, i+8]$ .
6. Any range  $A[l, r]$  can be decomposed into two ranges  $A[l, Y]$  and  $A[Z, r]$  where  $Y = l + 2x$  and  $Z = r - 2y$  such that  $Z \geq l$ ,  $Y \leq r$  and,  $A[l, Y]$ ,  $A[Z, r]$  overlap. Then,  $RMQ(A[i, j]) = \max(RMQ(A[l, Y]), RMQ(A[Z, r]))$
7. Total space cost  $O(n \log n)$ .

### Query Expansion

- User and documents may refer to a concept using different words
- Vocabulary mismatch
- Users often attempt to fix this problem manually
- Adding these synonyms should improve query performance

### Global Query Expansion

Retrieve synonyms from WordNet and Word2Vec.

### User relevance feedback

Relevance Feedback. User provides feedback to the search engine by indicating which results are relevant.

### **Pseudorelevance feedback**

- Take top-K results of original query
- Determine important/informative terms/topics (topic modelling!) shared by those documents Expand query by those terms
- No explicit user feedback needed (also called blind relevance feedback)

### **Indirect relevance feedback**

For a query look at what users click on in the result page

## **Lecture 5: Index Construction and Advanced Queries**

### **Static construction**

- Invert one batch
- Merge batches

### **Auxiliary Index**

- One static large static index on disk.
- As new documents arrive keep them in-memory in second index.

### **Logarithmic Index**

- Store index of size  $2^i \times n$
- Construction cost:  $N \log(N/n)$

### **Phrase queries**

- Inverted Index based (Positional Inverted Index)
- String matching indexes (Suffix Arrays)

### **More advanced queries**

- Wildcard/misspelling queries ( Sydney vs. Sidney: query  $S?dney$  )
- Regular expression queries ( “[J]ohn.\*@smith.com???” )
- Proximity queries (“president” close to “obama”)

## Lecture 6: IR Evaluation and re-ranking

### Hard to characterise the quality of a system's results \*

- a subjective problem, depends on the user's information need and how well the results meet that need.
- query is not the information need itself, but an expression thereof.

### Simplifying assumptions

- Retrieval is ad-hoc
- Effectiveness based on relevance

### recall is hard to calculate

### Precision-oriented metrics

- Precision@K: compute precision using only ranks 1 .. k
- Average Precision (AP): take average over precision@k for each k where rank k item is relevant; measure becomes rank sensitive
- Mean Average Precision (MAP): AP averaged across multiple queries

### Rank-biased precision

$$RBP = (1 - p) \times \sum_{i=1}^d r_i \times p^{i-1}$$

### Re-ranking

- Use BM25 as a first step in multi-stage retrieval system
- Use complex trained ranking model store rank the original BM25 ranking

### Rank objective

- Point-wise objective
- Pair-wise objective

## Lecture 7: Text Classification

### Text classification tasks

- Topic classification
- Sentiment analysis
- Authorship attribution
- Native-language identification
- Automatic fact-checking

### Building a Text classifier

- Identify a task of interest
- Collect an appropriate corpus
- Carry out annotation
- Select features
- Choose a machine learning algorithm
- Tune hyperparameters using held-out development data
- Repeat earlier steps as needed
- Train final model
- Evaluate model on held-out test data

### Naïve Bayes

- Finds the class with the highest likelihood under Bayes law.
- Naïvely assumes features are independent.
- Pros:
  - Fast to “train” and classify;
  - robust, low- variance;
  - good for low data situations;
  - optimal classifier if independence assumption is correct;
  - extremely simple to implement.
- Cons:
  - Independence assumption rarely holds;
  - low accuracy compared to similar methods in most situations;
  - smoothing required for unseen class/feature combinations

### Logistic Regression

- A linear model, but uses softmax “squashing” to get valid probability.
- Training maximizes probability of training data subject to regularization which encourages low or sparse weights.
- Pros:
  - A simple yet low-bias classifier;
  - unlike Naïve Bayes not confounded by diverse, correlated features
- Cons:
  - Slow to train;
  - Some feature scaling issues;
  - Choosing regularisation a nuisance but important since overfitting is a big problem

## Support vector machines

- Finds hyperplane which separates the training data with maximum margin.
- Pros:
  - Fast and accurate linear classifier;
  - Can do non-linearity with kernel trick;
  - Works well with huge feature sets
- Cons:
  - Multi-class classification awkward;
  - Feature scaling can be tricky;
  - Deals poorly with class imbalances;
  - Uninterpretable

## K-Nearest Neighbour

- Classify based on majority class of k-nearest training examples in feature space (Euclidean distance / Cosine distance)
- Pros:
  - Simple, effective;
  - No training required;
  - Inherently multiclass;
  - Optimal with infinite data
- Cons:
  - Have to select k;
  - Issues with unbalanced classes;
  - Often slow (need to find those k-neighbours);
  - Features must be selected carefully

## Decision tree

- Construct a tree where nodes correspond to tests on individual features.
- Pros:
  - In theory, very interpretable;
  - Fast to build and test;
  - Feature representation/scaling irrelevant;
  - Good for small feature sets, handles non-linearly-separable problems
- Cons:
  - In practice, often not that interpretable;
  - Highly redundant sub-trees;
  - Not competitive for large feature sets

## Random forests

- An ensemble classifier, Final class decision is majority vote of sub-classifiers
- Pros:
  - Usually more accurate and more robust than decision trees
  - training easily parallelised
- Cons:
  - Same negatives as decision trees
  - too slow with large feature sets



## Neural Networks

- An interconnected set of nodes typically arranged in layers.
- Pros:
  - Extremely powerful
  - State-of-the-art accuracy
- Cons:
  - Not an off-the-shelf classifier
  - Very difficult to choose good parameters
  - Slow to train
  - Prone to overfitting

## HyperParameter tuning

- Regularization hyperparameters penalize model complexity, used to prevent overfitting.
- For multiple hyperparameters, use grid search.

## Evaluation

- Accuracy = correct classifications/total classifications
- Precision =  $tp / (tp + fp)$
- Recall =  $tp / (tp + fn)$
- F1 =  $2 \text{ precision} * \text{recall} / (\text{precision} + \text{recall})$

## Lecture 8: N-gram language models

### Language models

Assign a probability to a sequence of words.

### Maximum Likelihood estimation

Estimate based on counts in our corpus

For unigram models,

$$P(w_i) = \frac{C(w_i)}{M}$$

For bigram models,

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}w_i)}{C(w_{i-1})}$$

For  $n$ -gram models generally,

$$P(w_i|w_{i-n+1} \dots w_{i-1}) = \frac{C(w_{i-n+1} \dots w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$

### Several problems

- Language has long distance effects (need a large  $n$ )
- Resulting probabilities are often very small (Use log probability to avoid numerical underflow)
- No probabilities for unseen words (Need to smooth the LM)

### Laplacian (Add-one) smoothing

For unigram models ( $V$ = the vocabulary),

$$P_{add1}(w_i) = \frac{C(w_i) + 1}{M + |V|}$$

For bigram models,

$$P_{add1}(w_i|w_{i-1}) = \frac{C(w_{i-1}w_i) + 1}{C(w_{i-1}) + |V|}$$

### Kneser-Ney smoothing

Back-off and Interpolation

$$P_{BO}(w_i|w_{i-2}, w_{i-1}) = \begin{cases} P^*(w_i|w_{i-2}, w_{i-1}) & \text{if } C(w_{i-2}, w_{i-1}, w_i) > 0 \\ \alpha(w_{i-2}, w_{i-1}) * P_{BO}(w_i|w_{i-1}) & \text{otherwise} \end{cases}$$
$$P_{interp}(w_i|w_{i-2}, w_{i-1}) = \lambda(w_{i-2}, w_{i-1}) P(w_i|w_{i-2}, w_{i-1}) + (1 - \lambda(w_{i-2}, w_{i-1})) P_{interp}(w_i|w_{i-1})$$

### Perplexity

$$PP(w_1, w_2, \dots, w_m) = \sqrt[m]{\frac{1}{P(w_1, w_2, \dots, w_m)}}$$

## Lecture 9: Lexical Semantics

### Lexical semantics

How the meanings of words connect to one another.

Manually constructed resources: lexicons, thesauri, ontologies, etc.

### Basic Lexical Relations

- Synonyms (same) and antonyms (opposite/complementary)
- Hypernyms (generic), hyponyms (specific)
- Holonyms (whole) and meronyms (part)

### Word similarity with paths

$$\text{simpath}(c_1, c_2) = 1/\text{pathlen}(c_1, c_2)$$

### Wu & Palmer similarity

$$\text{simwup}(c_1, c_2) = \frac{2 * \text{depth}(\text{LCS}(c_1, c_2))}{\text{depth}(c_1) + \text{depth}(c_2)}$$

### Lin Similarity

- \*  $P(c)$ : prob. that word in corpus is instance of concept  $c$

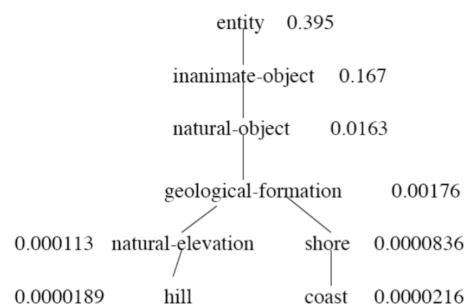
$$P(c) = \frac{\sum_{w \in \text{words}(c)} \text{count}(w)}{N}$$

- \* information content (IC)

$$IC(c) = -\log P(c)$$

- \* Lin distance

$$\text{simlin}(c_1, c_2) = \frac{2 * IC(\text{LCS}(c_1, c_2))}{IC(c_1) + IC(c_2)}$$



### Word sense disambiguation

- Supervised WSD
  - context is ambiguous
  - How big should context window be?
- Less supervised WSD
  - Choose sense whose dictionary gloss from WordNet most overlaps with the context.

Much modern work attempts to derive semantic information directly from corpora, without human intervention.

## Lecture 10: Distributional Semantics

### Lexical databases

- Manually constructed
  - Expensive
  - Human annotation can be biased and noisy
- Language is dynamic
  - New words: slang, terminology, etc.
  - New senses

### Distributional semantics

- Document co-occurrence often indicative of topic (document as context)
- Local context reflects a word's semantic class (word window as context)

### Two approaches:

- Count-based (Vector Space Models)
- Prediction-based

### Manipulating the VSM

- Weighting the values
- Creating low-dimensional dense vectors
- Comparing vectors

### Dimensionality reduction

- Singular value Decomposition ( $A = U \Sigma V$ )
- Latent semantic analysis (Truncating)

### Words as context

- Lists how often words appear with other words.
- The obvious problem with raw frequency: dominated by common words

### Pointwise mutual information

$$PMI(x, y) = \log_2 \frac{p(x, y)}{p(x)p(y)}$$

### Skip-gram: Factored Prediction

- Word embeddings should be similar to embeddings of neighbouring words
- Dissimilar to other words that don't occur nearby

## Lecture 11: Part of speech tagging

### POS Open classes

Nouns / Verbs / Adjectives / Adverbs

### POS Closed classes

Prepositions / Determiners / Pronouns / Conjunctions / Modals

### Automatic Taggers

- Rule-based taggers
- Statistical taggers
  - Unigram tagger
  - Classifier-based taggers
  - Hidden Markov Model (HMM) taggers

## Lecture 12: Neural sequence models

### FF-NN for Tagging

5 inputs: 3 x word embeddings and 2 x tag embeddings

1 output: vector of size  $|T|$ , using softmax

$$-\sum_i \log P(t_i | w_{i-2}, w_{i-1}, w_i, t_{i-2}, t_{i-1})$$

### Recurrent NNLMs

Tagging can benefit from context to left and right

#### Pros:

- Robust to word variation, typos, etc
- Excellent generalization, especially RNNs
- Flexible — forms the basis for many other models

#### Cons:

- Much slower than counts... but GPU acceleration
- Lots of classes (e.g., vocabulary)
- Not good for rare words... but pre-training on big corpora
- Data hungry, not so good on tiny data sets

## **Lecture 13: Information Extraction**

### **Machine learning in IE**

- Named Entity Recognition(NER):
  - sequence models such as seq. classifiers, HMMs or CRFs.
- Relation Extraction:
  - mostly classifiers, either binary or multi-class.

### **Dealing with adjacent entities: IOB tagging**

#### **Relation extraction**

- Fixed relation:
  - Rule-based
  - Supervised
  - Semi-supervised
  - Distant supervision
- Open relation:
  - Unsupervised
  - OpenIE

#### **Temporal expressions**

Anchoring: Information usually present in metadata.

Normalisation: mapping expressions to canonical forms.

#### **Event extraction**

Event ordering

## Lecture 14: Question Answering

### Definition:

Question Answering (“QA”) is the task of automatically determining the answer (set) for a natural language question.

### Question Processing

- Find key parts of question that will help retrieval.
- May reformulate question using templates.
- Predict expected answer type.

### Answer Extraction

Find a concise answer to the question, as a span in the text  
Framed as classification

### QA over structured KB

Natural language querying against knowledge bases using question parsing and logical inference.

## Lecture 15: Sequence Tagging: Hidden Markov Models

### HMMs for Tagging

Transition Matrix / Emission (observation) Matrix

### The Viterbi algorithm

- Complexity:  $O(T^2N)$ , where  $T$  is the size of the tag set and  $N$  is the length of the sequence.
- Because of the independence assumptions that decompose the problem (specifically, the Markov property).
- Good practice: work with log probabilities to prevent underflow (multiplications become sums).
- HMM is generative models.

## Lecture 16: Formal Language Theory & Finite State Automata

### What is a “language”?

a set of acceptable strings (e.g., sentences)

### Formal Language Theory

Formal apparatus to answer this question automatically, using a grammar.

### Key operations

- Membership
  - is the string part of the language?
- Scoring (requires weighting)
  - relax question to graded membership, how good an example of language is the string? (returning a number)
- Transduction
  - input one string, output another
  - A form of translation, but used extensively e.g., tagging = translating from words to tags

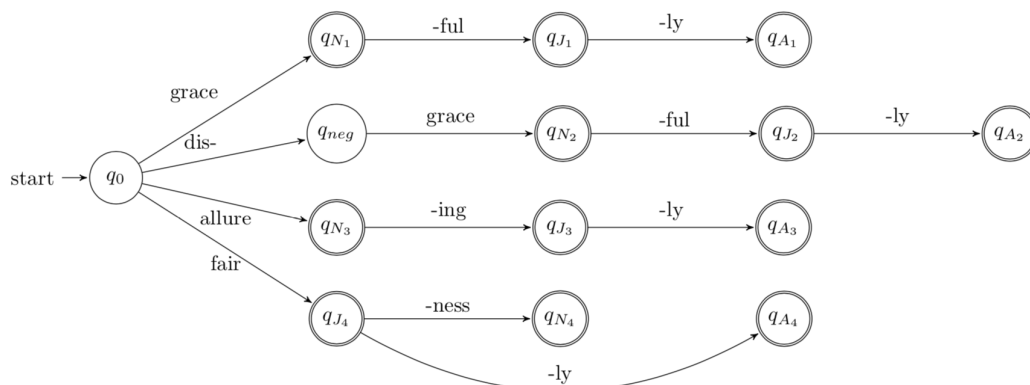
### Accepted by regular expression which supports the following operations:

- Symbol drawn from alphabet,  $\Sigma$
- Empty string,  $\epsilon$
- Concatenation of two regular expressions, RS
- Alternation of two regular expressions, R|S
- Kleene star for 0 or more repeats,  $R^*$
- Parenthesis () to define scope of operations

### Finite State Acceptors

Accepts strings if there is path from  $q_0$  to a final state with transitions matching each symbol.

### FSA for word morphology





## Lecture 17: Context-free Grammars

### Basics of Context-free grammars

- Symbols
  - Terminal: word such as book
  - Non-terminal: syntactic label such as NP or NN
  - Convention to use upper and lower-case to distinguish, or else “quotes” for terminals
- Productions (rules)  $W \rightarrow XYZ$ 
  - Exactly one non-terminal on left-hand side (LHS)
  - An ordered list of symbols on right-hand side (RHS) can be Terminals or Non-terminals

### Regular expressions as CFGs

e.g.  $[A-Z][a-z]^*$   
 $S \rightarrow U \quad S \rightarrow ULS$   
 $U \rightarrow "A" \quad U \rightarrow "B" \quad \dots \quad U \rightarrow "Z"$   
 $LS \rightarrow L \quad LS \rightarrow LLS$   
 $L \rightarrow "a" \quad L \rightarrow "b" \quad \dots \quad L \rightarrow "z"$

**The class of regular languages is a subset of the context-free languages, which are specified using a CFG.**

### CFGs vs regular grammars

- Regular grammars
  - describe a smaller class of languages
  - can be parsed using finite state machines (FSA, FST)
- CFGs
  - can describe hierarchical groupings
  - requires more complex automata to parse (PDA)

### CFG trees

- Generation corresponds to a syntactic tree
- Non-terminals are internal nodes
- Terminals are leaves
- Often more than one tree can describe a string

### Parsing CFGs

- Bottom-up
  - Start with words, work up towards S
  - CYK parsing
- Top-down
  - Start with S, work down towards words
  - Earley parsing (not covered)

## **Lecture 18: Probabilistic Parsing**

### **Basics of Probabilistic CFGs**

- As for CFGs, same symbol set and same productions.
- In addition, store a probability with each production.
- Probability values denote conditional.
- Each probability must be positive values, between 0 and 1, and the sum must to be 1.
- Grammar / Lexicon

### **Resolving parse ambiguity**

Get the multiplication of all elements in trees.

S in the top-right corner of parse table indicates success

## Lecture 19: Dependency Grammar & Parsing

### Dependency G vs. Phrase-Structure G

- phrase-structure grammars assume a constituency tree which identifies the phrases in a sentence. Based on idea that these phrases are interchangeable (e.g., swap an NP for another NP) and maintain grammaticality.
- Dependency grammar offers a simpler approach: describe binary relations between pairs of words. Namely, between heads and dependents.

### What is a Dependency?

- Links between a head word and its dependent words in the sentence: either syntactic roles or modifier relations.
- Dependency tree more directly represents the core of the sentence: who did what to whom?

### Dependency tree

- Dependency edges form a tree
  - each node is a word token
  - one node is chosen as the root
  - directed edges link heads and their dependents
- Cf. phrase-structure grammars
  - forms a hierarchical tree
  - word tokens are the leaves
  - internal nodes are 'constituent phrases' e.g., NP, VP etc
- Both use part-of-speech

### Projectivity

- A tree is projective if, for all arcs from head to dependent.
- There is a path from head to words that lies between the head and the dependent.
- The tree can be drawn on a plane without any arcs crossing.

### Dependency grammar

- In sense of generative grammar.
- Cannot be said to define a language, unlike a context free grammar.
- Any structure is valid, job of probabilistic model to differentiate between poor and good alternatives.
- Many more phrase-structure treebanks, which can be converted into dependencies.


### Dependency parsing

- Task of finding the best structure for a given input sentence.
- Graph-based: uses chart over possible parses, and dynamic programming to solve for the maximum.
- Transition-based: treats problem as incremental sequence of decisions over next action in a state machine.

## Transition based parsing

- Maintain two data structures
  - buffer = input words yet to be processed
  - stack = head words currently being processed
- Two types of transitions
  - shift = move word from buffer on to top of stack
  - arc = add arc (left/right) between top two items on stack (and remove dependent from stack)
- Always results in a projective tree.


Buffer	Stack	Action
I shot an elephant in my pyjamas		Shift
shot an elephant in my pyjamas	I	Shift
an elephant in my pyjamas	I, shot	Arc-left
an elephant in my pyjamas	shot	Shift
elephant in my pyjamas	shot, an	Shift
in my pyjamas	shot, an, elephant	Arc-left
in my pyjamas	shot, elephant	Arc-right
in my pyjamas	shot	Shift
...	...	...
	shot	<done>



- **How do we know when to arc and whether to add left or right facing arcs?**  
 Uses an “oracle” sequence of parser actions. Predict next action in sequence, and update when model disagrees with gold action.

## Graph based parsing

- Can consider as a CFG, where lexical items (heads) are non-terminals.
- Score of parse assumed to decompose into pairwise dependencies.
- production shot → shot in means arc-right from “shot” to “in”.



## **Lecture 20: Discourse**

### **Discourse**

a coherent, structured group of sentences (utterances)

### **Discourse segmentation**

Assumption: text can be divided into a number of discrete, contiguous sections.

Task: classifying whether a boundary exists between any two sentences.

### **An unsupervised approach (Text Tiling)**

looking for points of low lexical cohesion.

### **Supervised discourse segmentation**

Apply a binary classifier to identify boundaries.

- distributional semantics
- coreference cues
- discourse markers

### **Discourse parsing**

- A proper discourse must be coherent
- Discourse units (DUs) are related by specific coherence relations
- Two related DUs form a new DUs
- All DUs in a coherent discourse must be related
- A discourse will form a tree, which can be parsed

### **Anaphors**

linguistic expressions that refer back to earlier elements in the text

### **Antecedent Restrictions**

- Pronouns must agree in number with their antecedents
- Pronouns must agree in gender with their antecedents
- Pronouns whose antecedents are the subject of the same syntactic clause must be reflexive (...self)

### **The Centering Algorithm**

at any given moment, discourse is focused on a single entity, the “center”.

## **Lecture 21: Machine translation: word-based models**

### **Noisy channel MT**

- $\hat{e} = \operatorname{argmax}_e P(e) P(f|e)$
- $P(f|e)$  rewards good translations, but permissive of disfluent  $e$
- $P(e)$  rewards  $e$  which look like fluent English, and helps put words in the correct order.

### **How to learn the LM and TM**

LM: based on text frequencies in large monolingual corpora (as seen in previous lecture)

TM: based on word co-occurrences in parallel texts

### **maximum likelihood estimator**

## **Lecture 22: Machine translation: phrase based & Neural Encoder- decoder**

### **Phrase based MT**

Treats n-grams as translation units, referred to as 'phrases' (not linguistic phrases, just adjacent words)

### **Finding & scoring phrase pairs**

- "Extract" phrase pairs as contiguous chunks in word aligned text;
- Compute counts over the whole corpus;
- Normalise counts to produce 'probabilities';

### **Neural Machine translation**

sequence 2 sequence (encoder and decoder)

MT Evaluation: BLEU